

Driving Continuous Improvement with Service Levels



Elton Stoneman

Consultant & Trainer

@EltonStoneman blog.sixeyed.com


Service Level Objectives



Response time

99.9% of requests within 500ms

14 days = 20,160 minutes | Error Budget = 20 minutes



- Baseline performance borderline
- Release degrades performance

Lessons Learned


- ✓ [] [] []
- ✓ [] [] [] []
- ✗ [] [] [] []
- ✗ [] []

Action Items

- [] [] []
- [] []
- [] [] []



Service Levels



SLI
Service Level
Indicator



SLO
Service Level
Objective



SLA
Service Level
Agreement



SLI

X is true


SLO

for Y period of time

SLA

... or else

Matthew Flaming, New Relic



Designing SLOs



Customer journeys



Improvable targets




Measure breaches



A Globomantics Company



Amila




Scenario: Designing SLOs and SLIs




Amila

What are the most important user journeys?





- Event search
- Event details
- Reserve tickets
- Buy tickets



Prioritized User Journeys



#1

Buy tickets



#2

Event search



#3

Event details




#4

Reserve tickets





Prioritized SLOs

		<i>Customer Impact</i>	<i>Business Impact</i>
#1	Buy tickets	★ ★ ★	★ ★ ★
#2	Event search	★ ★ ★	★ ★
#3	Event details	★ ★	★
#4	Reserve tickets	★	★




What is a
"good"
buy-ticket
experience?






- Correctly and quickly
- Latency & status tracked
- Histogram for latency








- 95th percentile
- Service performance
- Client performance






What SLO
is
achievable
and
desirable?






- Data-driven targets
- 99.9% success rate
- 500ms latency





- Biz/dev alignment
- Shorter for volatile
- Longer for stable



99.9% of ticket purchases
complete **successfully**
within **500ms**
over a **one-month** period

Understanding the Goals of SLOs

Reactive SLOs



**"Just about"
good enough**



**Minimizing
intervention**



**Hiding
underlying issues**

Development Team

SRE Team



New release ready!



*No way! You've used
90% of your error budget*

Reactive SLOs



Response time

99% of home page requests within 900ms

14 days = 20,160 minutes | Error Budget = 201 minutes



No breaches for 427 weeks!


Proactive SLOs



**Positive customer
experience**



**Error-budget
tracking**




**Driving
improvement**

Proactive SLOs






Ticket purchase 99.9% successful within 550ms

1 month \approx 43,200 minutes | Error Budget \approx 43 minutes






Breaches decrease

Scenario: Evolving SLOs Over Time






Can the
product
devs
measure
the SLO?







- End-to-end testing
- Includes SLO journey
- Not timed




- Include timings
- Collect baseline
- Add SLO test







- Predictable latency
- Synthetic tests
- Correlate to prod







What type of action items from an SLO breach?





- Big issue -> postmortem
- Code or rollout issue
- Product or process actions




- Error budget policy
- Postmortems:
 - 25% in one issue
 - Budget exhausted




Error Budget Policy




Ticket purchase 99.9% successful within 550ms



- **Code issue**




- **Deployment issue**



What if the
SLO is
never
breached?






- Not ambitious enough
- No failures
- No improvement

SLO Reviews



Ticket purchase

99.9% successful within 550ms



Error budget not consumed



- **Reduce latency target**
- **Improve success target**
- **Measure at 99th percentile**

Using SLOs in Feedback Loops

Attainable and Desirable SLOs



Ticket purchase

successful within

- **99.9% @ 500ms**
- Good for customers
- Achievable for SRE team

- **99.5% @ 800ms**
- OK for customers
- Easy for SRE team

Running at SLO Performance



Reduce scale



Synthetic outage




Monitor reaction


- End-to-end test
- Timed at 1.6x prod




- Synthetic test
- Times full experience




dev



prod



Timeline




Lessons Learned



Action Items


-
-
-





Scenario: Monitoring SRE Performance



How do you
monitor the
performance
of the SRE
team?






...




How can you drive improvement from this data?






- Monitor incidents
- Record MTBF and MTTR
- Trend product & process



MTTR




- Repair
- Recovery
- Respond
- Resolve





- SRE backlog
- Tied to SLOs
- Otherwise?




Module Summary

Measuring the Measurements



**Application
performance**



**SLO breaches
and recovery**



**Incident
management**

Product Backlog




Safe restarts

3 days

Dev

1

x



Health checks

2 days

SRE

2

y



Caching

5 days

SRE

3

z

Summary



Designing service levels

- Customer-focused SLOs
- Meaningful SLIs
- Achievable and desirable

Evolving SLOs over time

- End-to-end and synthetic tests
- Review and tighten goals


Driving improvement with SLOs

- Monitoring MTBF and MTTR
- Product feedback loops


Course Layout



Architecture



Observability



Improvement



Technical Deep Dives

Prometheus, Jenkins, Kubernetes...

Elton Stoneman

We're Done!



So...

- **Please leave a rating**
- **Follow @EltonStoneman on Twitter**
- **And watch my other courses 😊**