

Implementing SRE Best Practices



Karun Subramanian
IT Operations Expert

@karunops www.karunsubramanian.com



Overview



Utilizing the circuit breaker pattern

**Comparing DNS load balancing vs.
dedicated load balancer**

Transitioning to canary-based deployment

Understanding distributed consensus

Designing auto scaling based on load

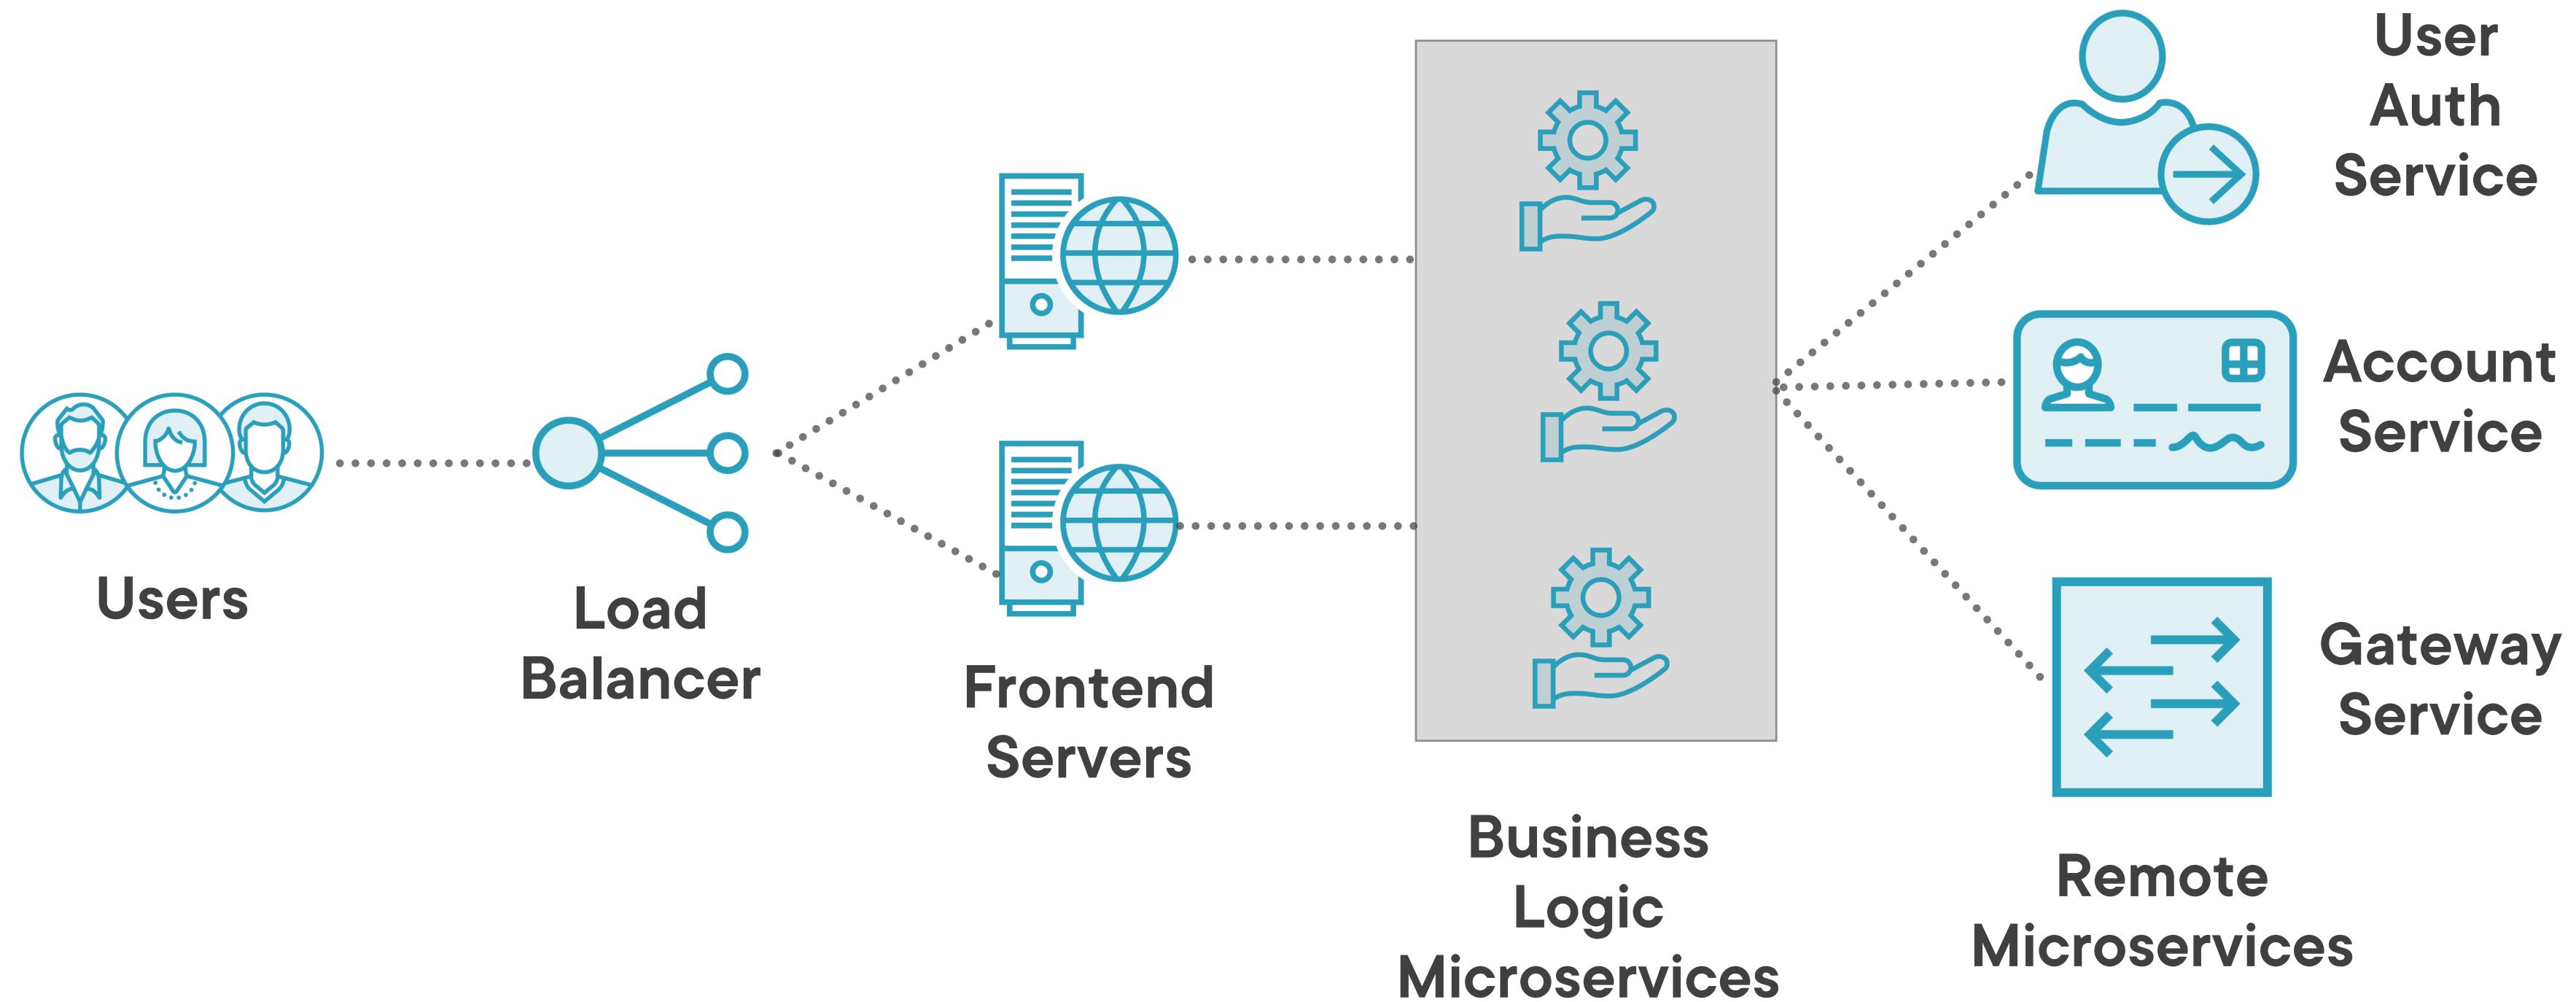
Implementing effective health checks



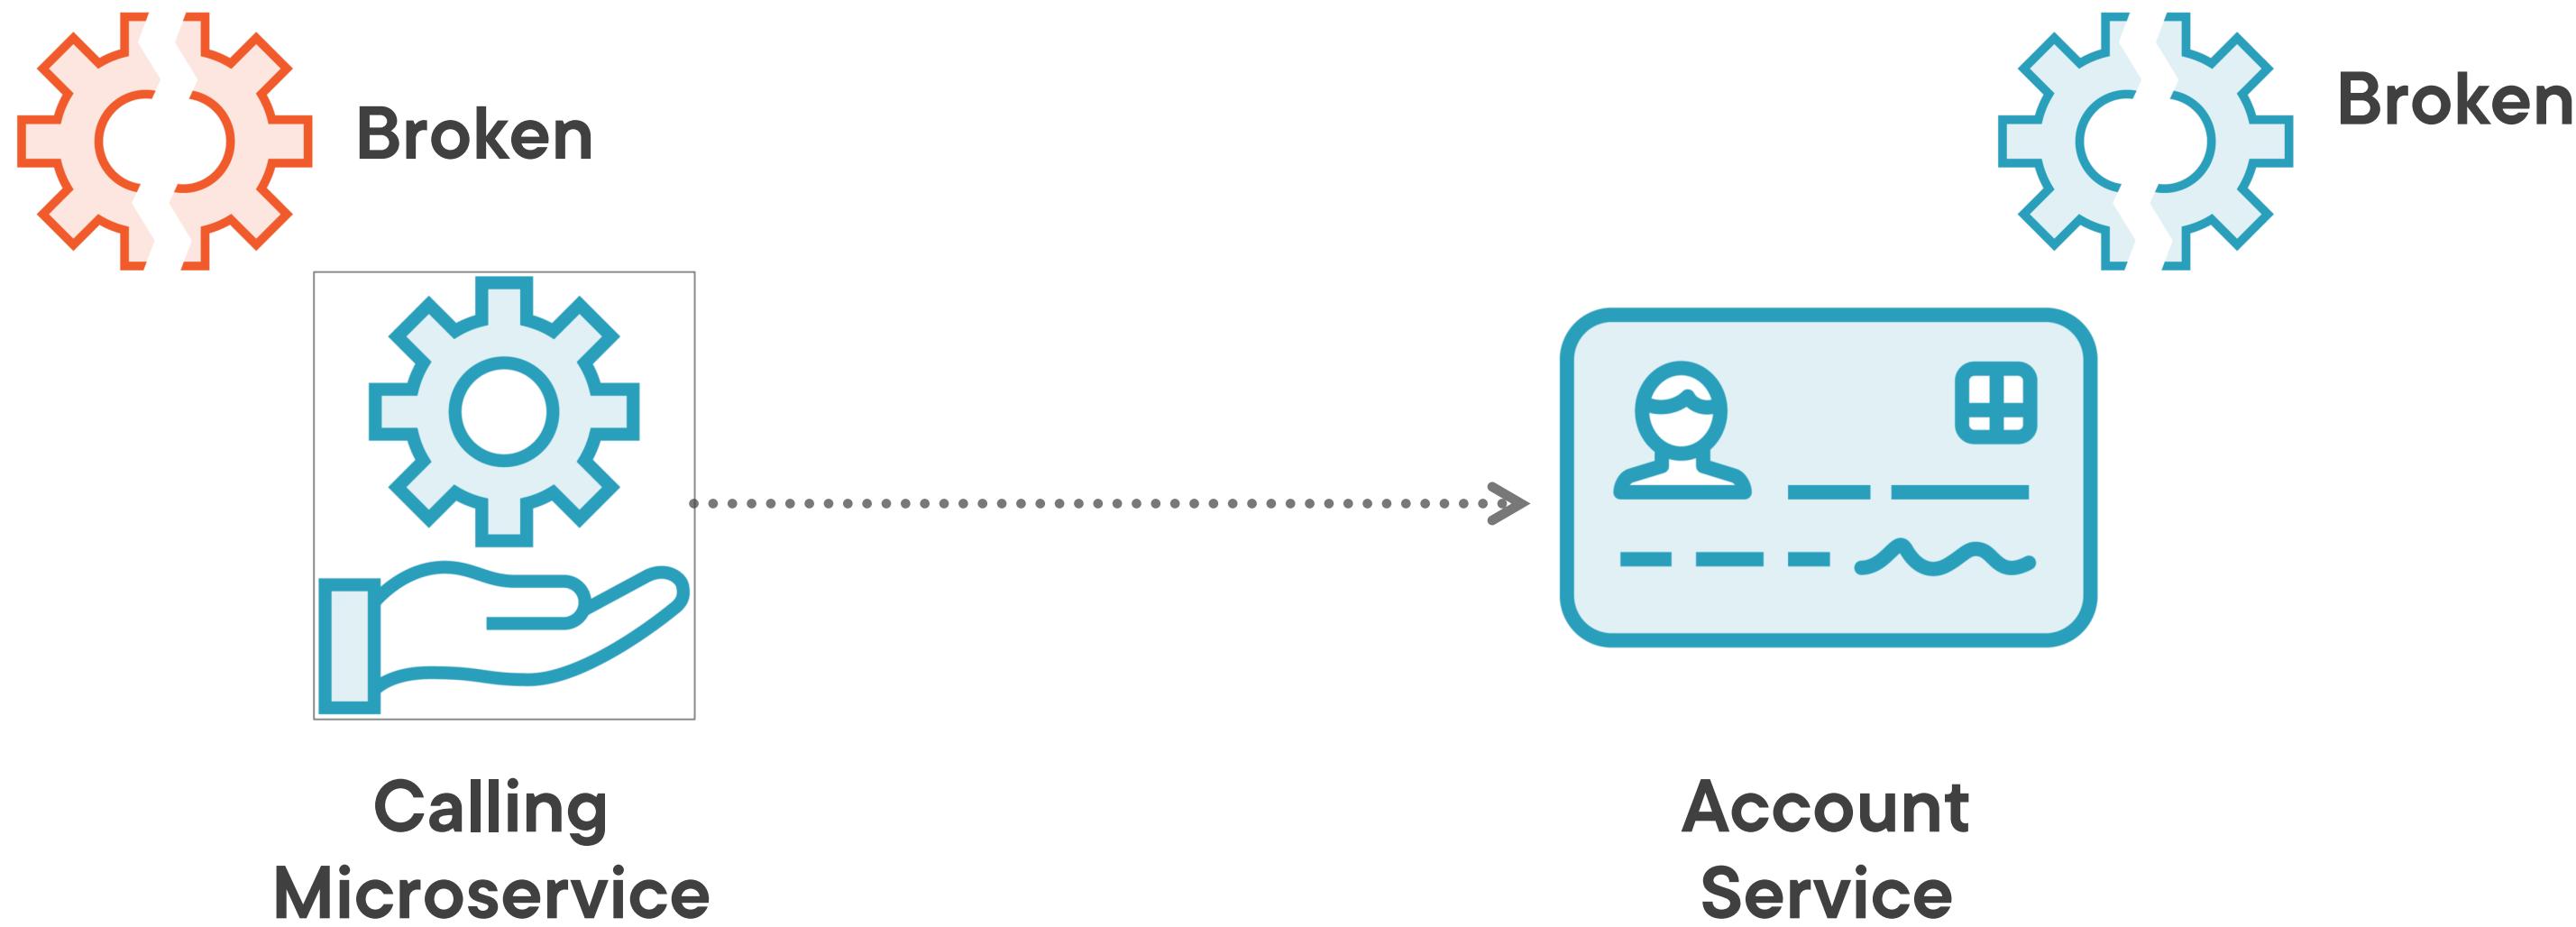
Utilizing the Circuit Breaker Pattern



Typical Distributed Architecture



Deadly Pitfall



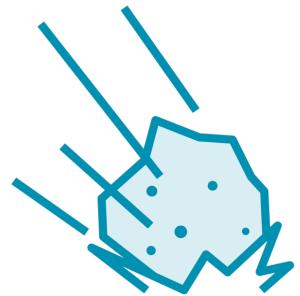
**Poorly performing remote web services can cause
the calling services to break**



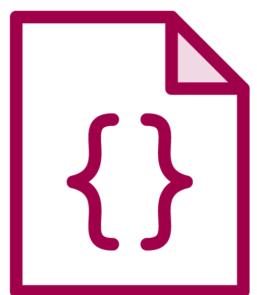
Circuit Breaker Pattern



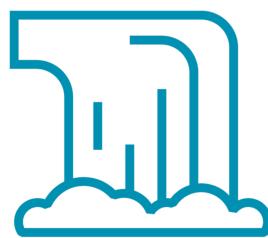
When a remote service fails, fail fast



Helps to deal with failures to minimize user impact (failing gracefully)



Instead of bringing down the entire application, helps to run the application with reduced functionality



Avoids cascading failures



Circuit Breaker Logic



Track the number of failures encountered while calling a remote service



Fail (open the circuit) when a pre-defined count is reached



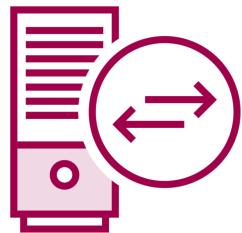
Wait for a pre-defined time and retry connecting to the remote service



Upon successful connection, close the circuit. Upon failure, restart the timer



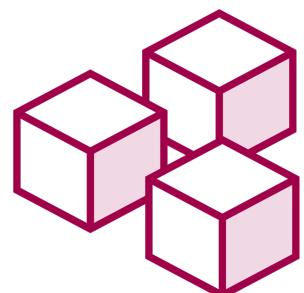
Implementing Circuit Breaker



Use a proxy layer such as istio



Use Hystrix, opensource library from Netflix



Istio is technology independent (Blackbox approach) and sits outside of your application



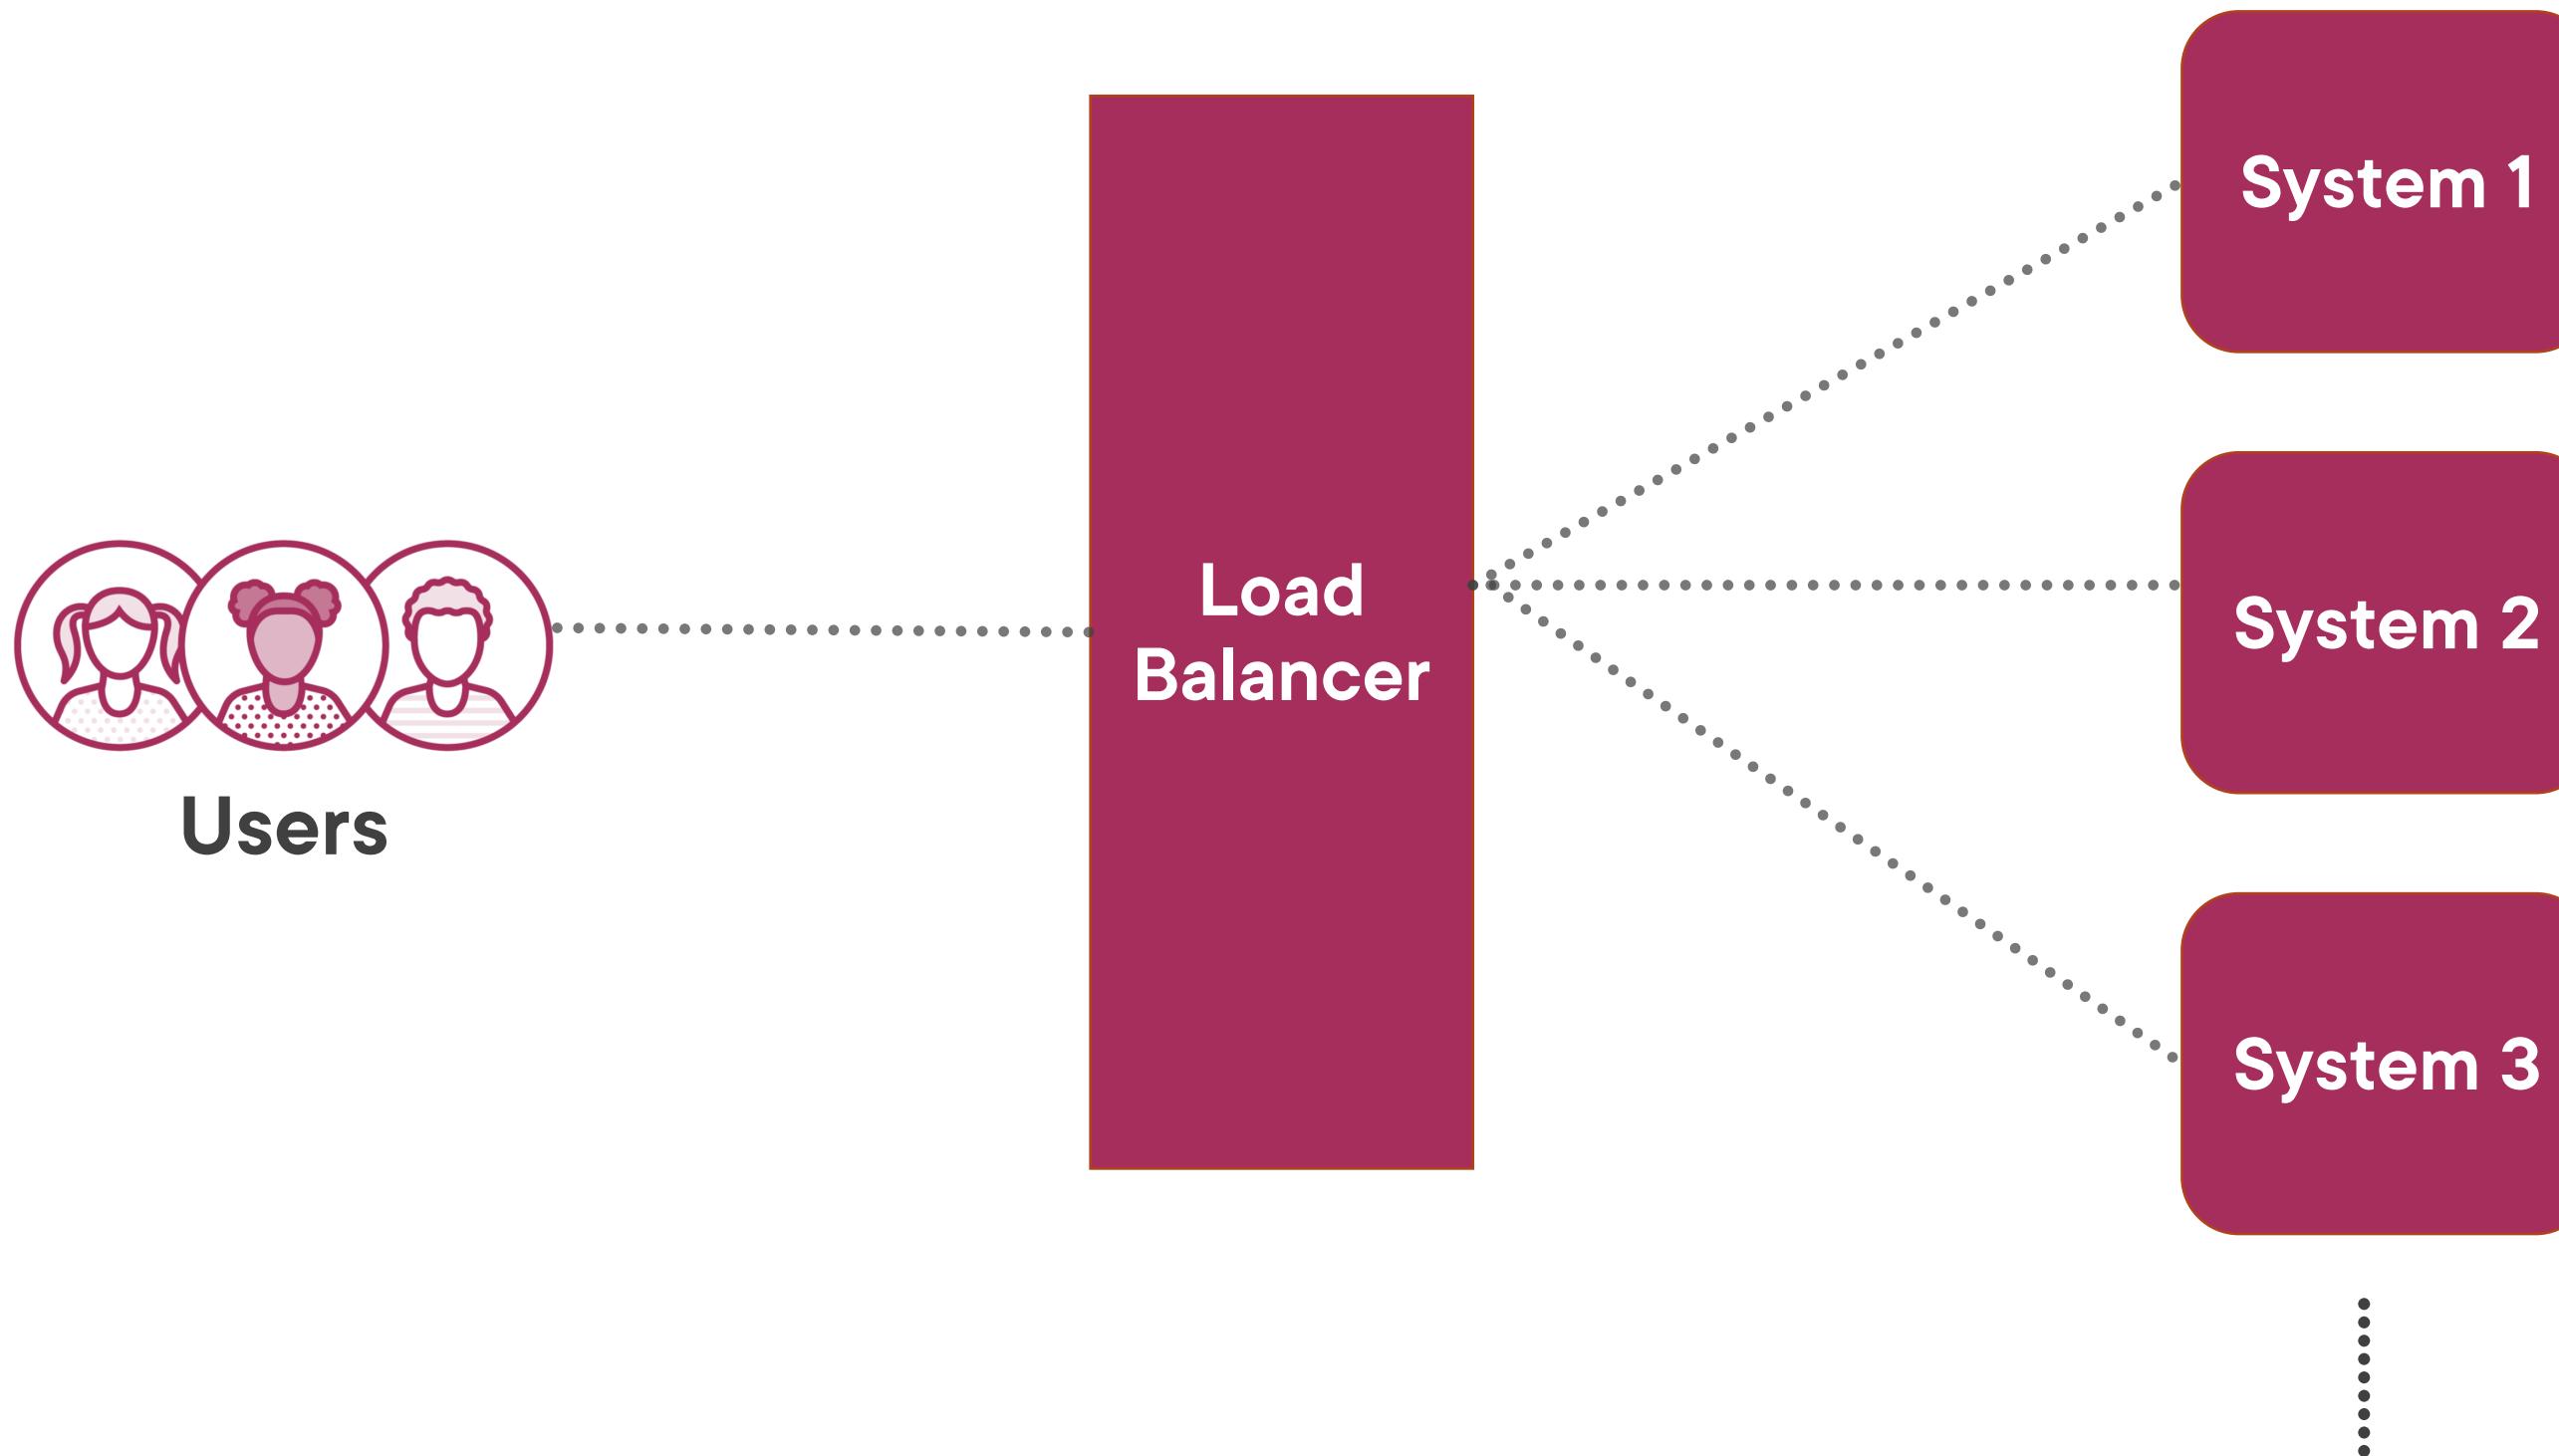
Hystrix is built with your code (Whitebox approach) but may provide additional functionalities



Implementing Effective Load Balancing



Load Balancers



Two Functions of Load Balancing

High Availability

Failures can be tolerated in a pool of servers

Performance

By adding more servers in a pool, effective performance is increased by horizontally spreading the load



Three Basic Ways to Load Balance

DNS

Closest to the client.
Multiple A records
enable client to
choose one at random.

Hardware Load Balancer

Feature-rich. Highly
performant. Typically
used within
datacenter.

Software Load Balancer

Relies on host
hardware. Can be
useful for dev
environments.



DNS Load Balancing

DNS server returns multiple IP addresses when a client looks up a hostname

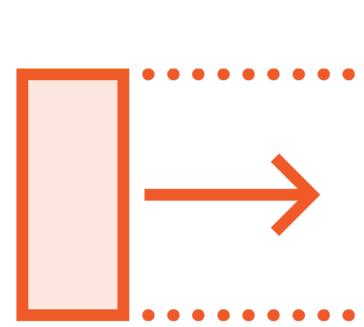
DNS server can be configured to choose a datacenter closest to the client's geographical location

DNS server can also determine the targets based on load/capacity of the datacenter

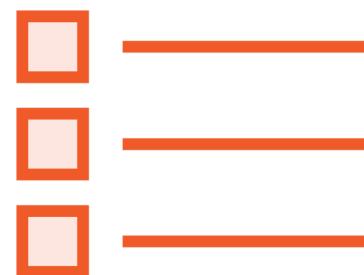
Among the IP addresses returned, client chooses an IP at random



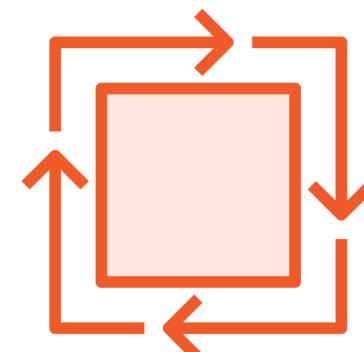
Potential Problems with DNS Load Balancing



Not easy to control (client picks an IP randomly)



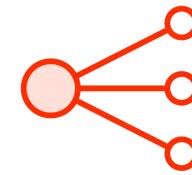
Requests can end up being sent to a server that is down



Clients cache DNS entries heavily and it is not easy to clear them



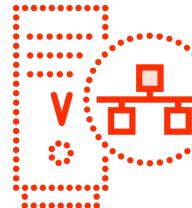
Dedicated Load Balancers



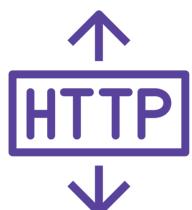
Layer3 (Network), Layer4 (Transport) and Layer7(Application) load balancers



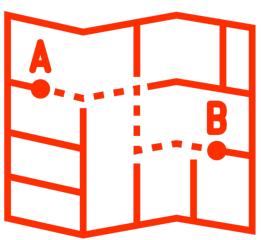
L3 load balancers operate on source and destination IP addresses



L4 load balancers operate on IP + Port (TCP protocol)



L7 load balancers operate by processing application data such as HTTP URL, cookie, header etc.



L7 load balancers provide rich functionality for routing



Load Balancing Methods

Round Robin

Targets are rotated in a loop

Weighted Round Robin

Manually assigned weight influences the backend target

Least Loaded

Based on the reported load by the backend

Least Loaded with Slow Start

Prevents flooding the least loaded backend

Utilization Limit

Based on the reported utilization (QPS) by the backend



Hardware vs. Software Load Balancers

Hardware Based Load Balancers

Purpose built appliances

Feature rich

Can be expensive

Examples: F5, Imperva

Software Based Load Balancers

Software running on commodity hardware

Comparable. Fast evolving

Relatively less expensive

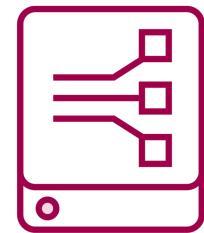
Examples: HAProxy, NGINX



Transitioning to Canary-based Deployment



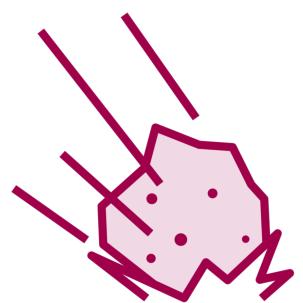
What is Canary-based Deployment?



Rolling out changes to small number of targets before releasing to all



Like an insurance policy against bad releases



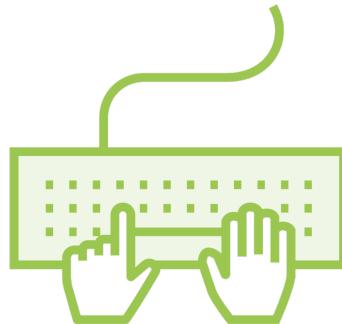
Can help reduce or eliminate user impact



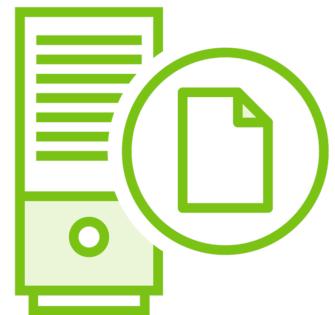
Not a substitute for integration/load tests



How to Perform Canary-based Deployment?



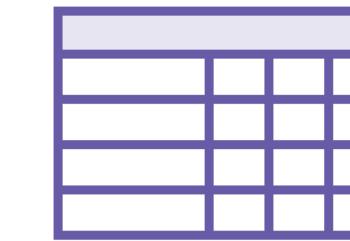
Exact method will depend on your environment



Release on one server first



Ensure that there are no catastrophic failures (server crash)



Release up to 1% of the servers



Proceed to release for early adaptors



Release for all users



A failed canary is a serious issue. Analyze and implement robust testing process so that issues can be caught during testing.



Understanding Distributed Consensus



Inherent Problems with Distributed Systems

Hardware fails

Nodes in a distributed system can randomly fail

Network outages are inevitable

We cannot always guarantee 100% network connectivity

Need for a consistent view

All nodes in a distributed system must achieve a consistent view of system state

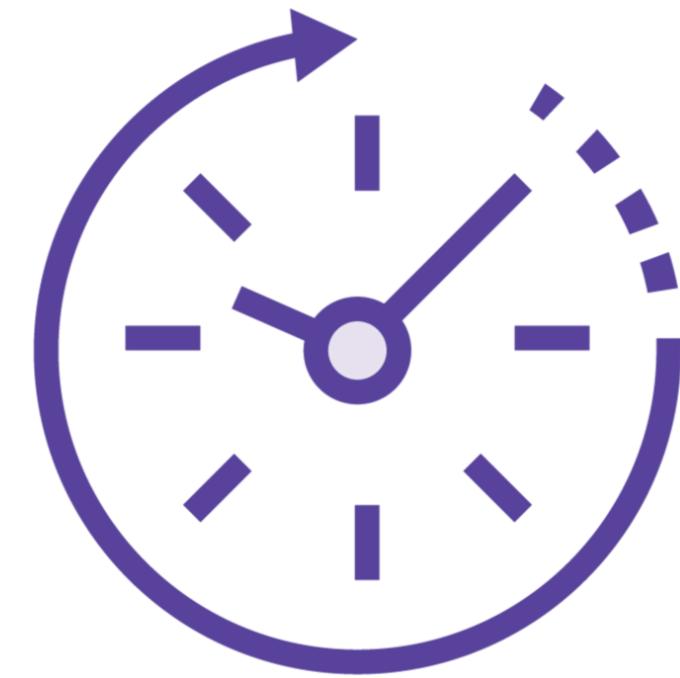


CAP Theorem

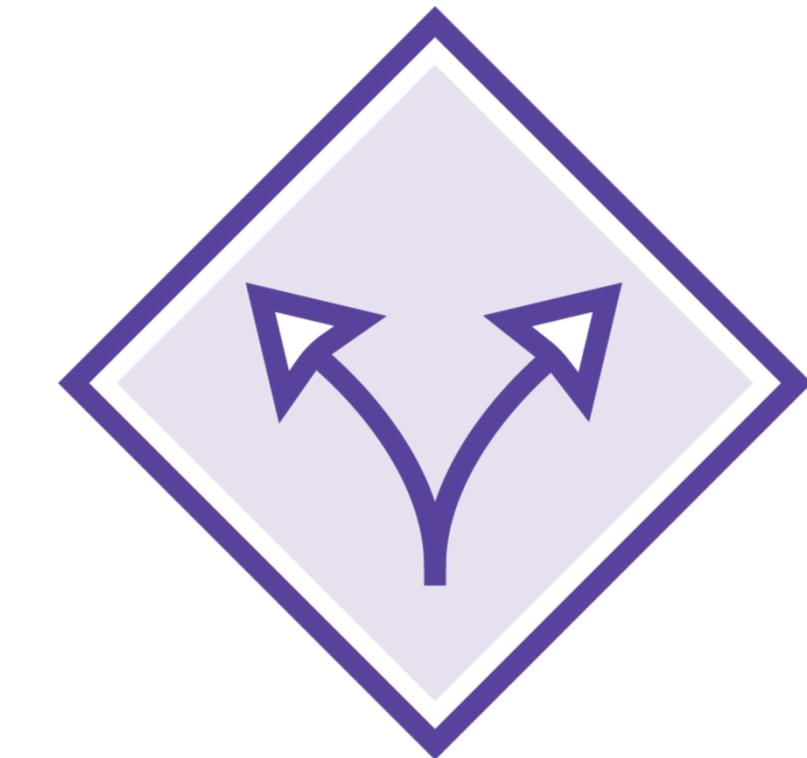
A distributed system cannot simultaneously have these properties



Consistency
Consistent views of
data at each node



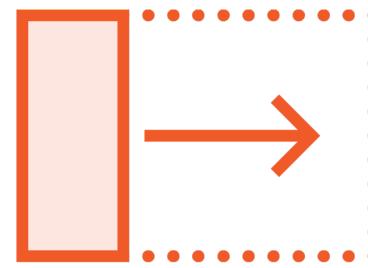
Availability
Availability of data at
each node



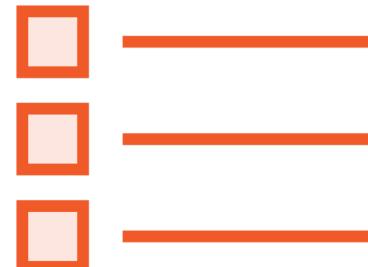
Network Partition
Tolerance to Network
failures (which results
in network partitions)



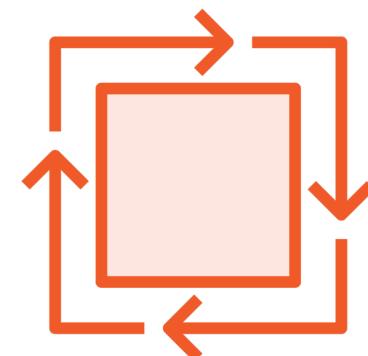
Distributed Consensus



Protocol to build reliable distributed systems



Ad-hoc methods such as heartbeats are not sufficient



Over the years, several protocols such as Paxos, Raft and Zab have been developed



Distributed Consensus Algorithm (Paxos)

Paxos was the original solution to the distributed consensus problem

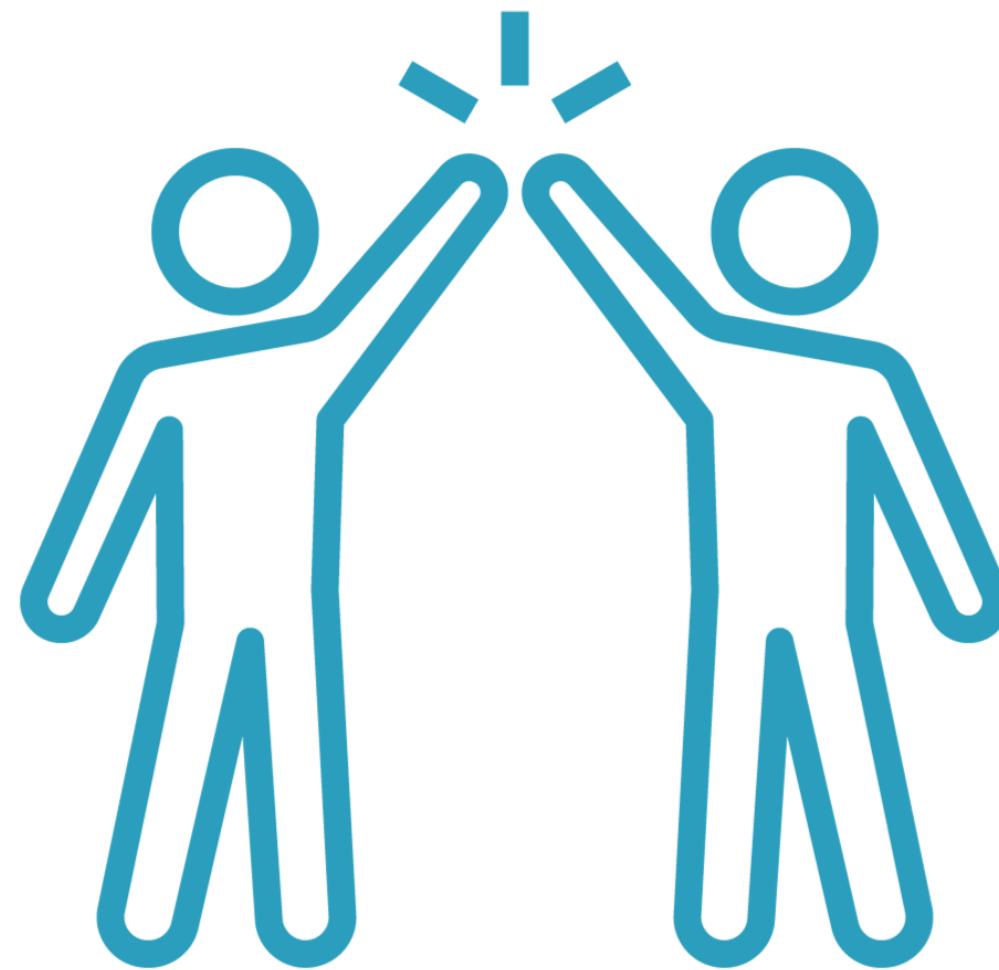
Systems in a distributed system send a series of proposals with a unique sequence number

When majority of the processes accept the proposal, the sender sends a commit message

The strict sequence of proposals solves the problem of ordering



Distributed Consensus Implementations



In addition to Paxos, other algorithms are available

- Raft
- Zab
- Mencius

Distributed consensus is generally implemented using commercial/opensource services

- Zookeeper
- Consul
- Etcd



Designing Auto Scaling



Autoscaling is a process by which the number of servers in a server farm is automatically increased or decreased based on the load



Four Benefits of Autoscaling

Reduces cost by running only the required servers

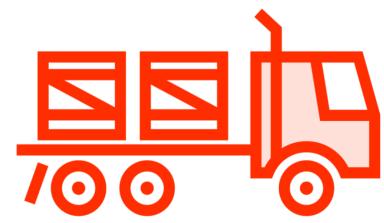
Flexibility to run less time-sensitive workload during low traffic

Automatically replace unhealthy servers (most commercial cloud vendors)

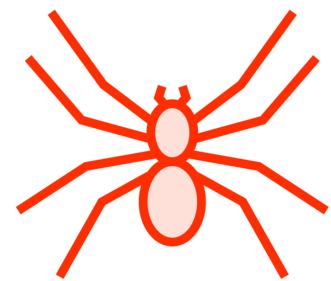
Increased reliability and uptime



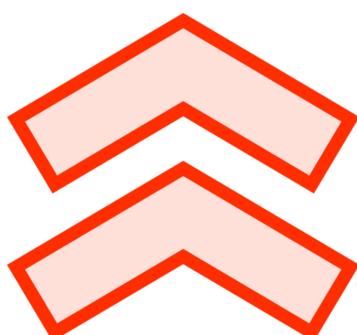
Problems with Autoscaling



A dependent backend can get overwhelmed



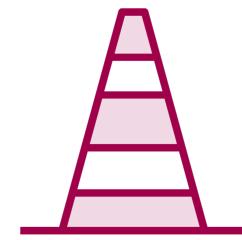
Software bugs can trigger autoscaler to expand the server farm abruptly



Load balancing may not be intelligent enough to consider newly added servers (i.e. warm up period)



Autoscaling Best Practices



Scaling down is more sensitive/dangerous than scaling up. Fully test the scale-down scenarios



Ensure the backend systems (database, remote web service, etc.) can handle increased load



Configure upper limit on the number of servers



Have a kill switch that can be easily used



Three Systems to Act in Concert

Load Balancing

Minimize latency by
routing location
closest to the user

Nginx/HAProxy

Load Shedding

Serves requests that
we can and drop
excess traffic

Zuul/Envoy

Autoscaling

Based on load,
automatically scale up
or down

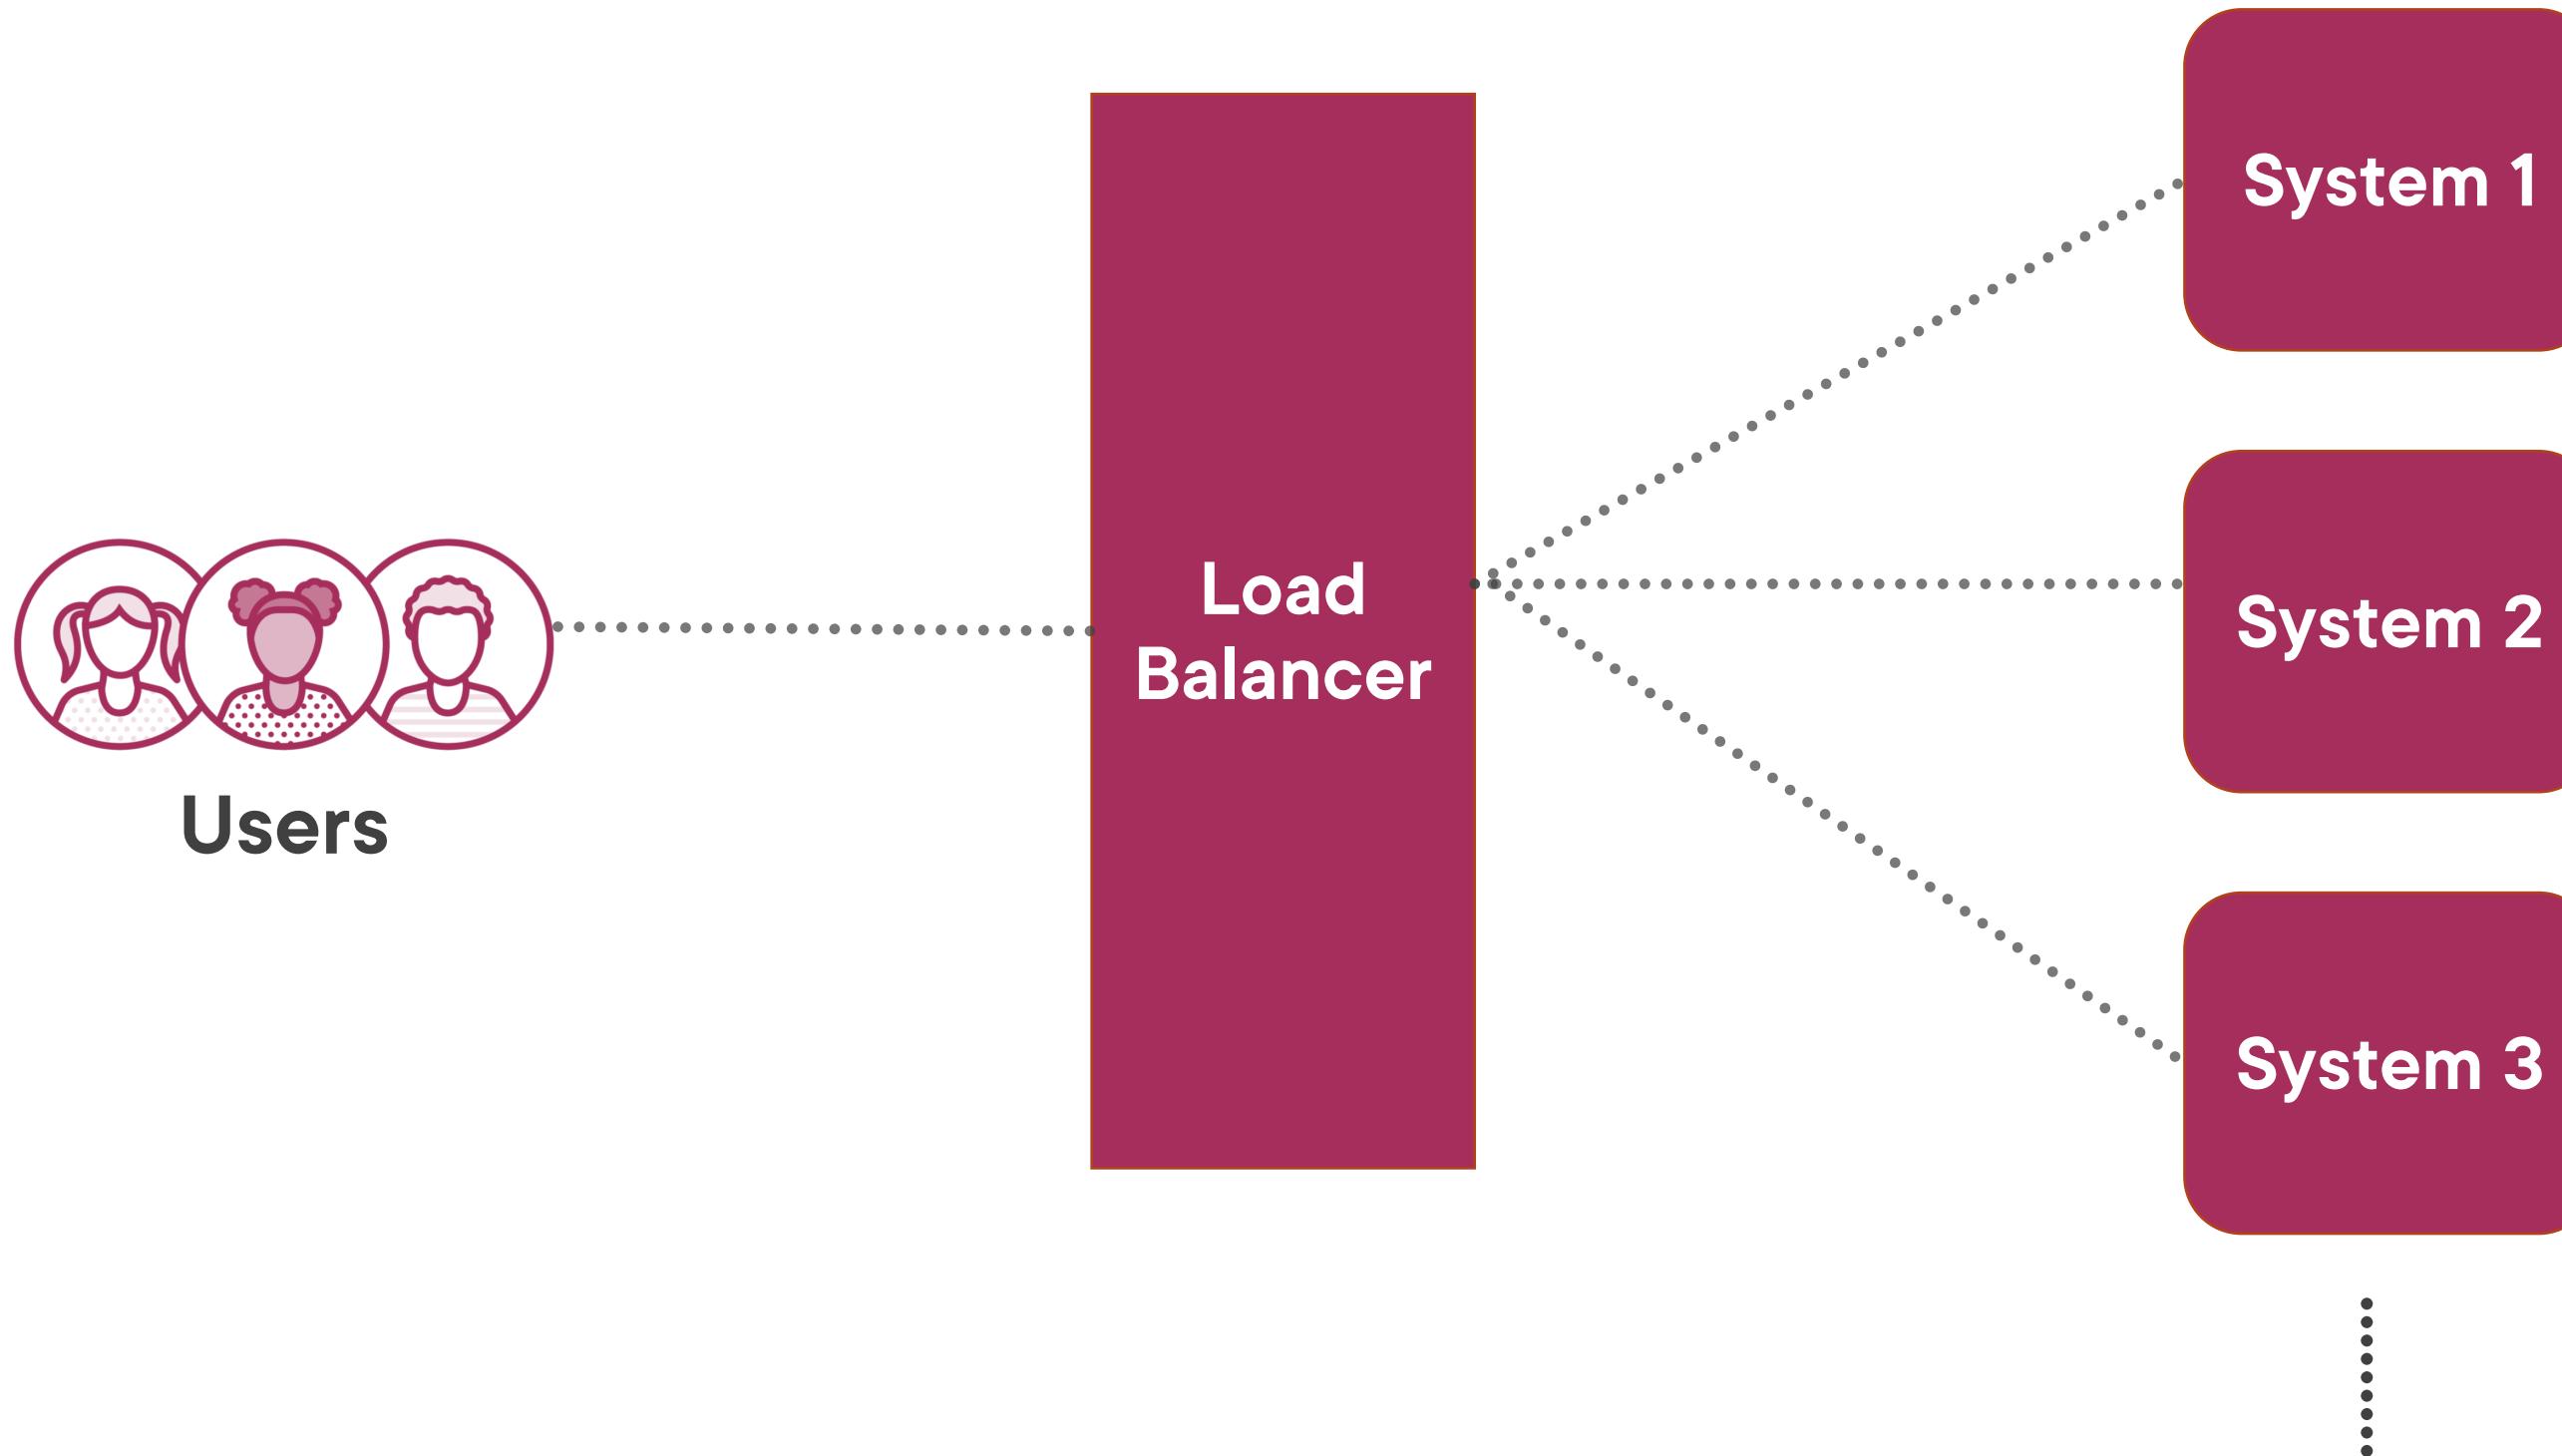
AWS/GCP/Azure



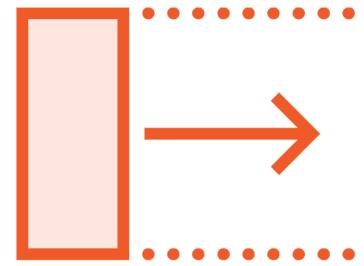
Implementing Effective Health Checks



Load Balancing Revisited



Load Balancers



Load balancers direct traffic to a set of backend servers



Load balancers need to know which servers are alive and healthy



Load balancers use health checks to determine the healthy servers



Effective Health Checks

Simple

Monitors for the availability of the backend server

Content verification

Sends a request and examines the response from the backend server

Reliable health checks

Health checks must cover situations such as ‘hung’ backend servers



Health checks with sophisticated content verification can increase network traffic. It can also be slow to mark a pool member as ‘down’



Summary



Utilize the circuit breaker pattern to fast-fail

- Hystrix (Whitebox approach)
- Istio (Blackbox approach)

Design load balancing with a mix of DNS and dedicated load balancers

- DNS load balancing may not be reliable

You must use canary releases

- Canary is not a replacement for testing

CAP theorem states that you cannot simultaneously have

- Consistency
- Availability
- Tolerance to network partition



Summary



When designing auto scaling,

- Ensure backend systems can handle load
- Have a kill switch
- Develop accurate procedures for scale-down

Implement load balancing, load shedding and autoscaling to work together

Configure reliable load balancer health checks

- Simple ping is fast but may not be reliable
- Content-based is reliable but it may increase network bandwidth usage



Up Next:
Benefits of SRE

